Dissertation
Towards a Framework for DHT Distributed Computing

Andrew Rosen

Georgia State University

May 13th, 2016

Georgia State
University

Table of Contents

Georgia State
University

Objective

Our objective is to create a generalized framework for distributed computing using Distributed Hash Tables.

Georgia State University

Objective

Our objective is to create a generalized framework for distributed computing using Distributed Hash Tables.

Or

Georgia State University

Objective

Our objective is to create a generalized framework for distributed computing using Distributed Hash Tables.

Or

We want to build a completely decentralized distributed computing framework.

Georgia State University

Introduction Background Previous Work   UrDHT Autonomous Load-Balancing   Conclusion
○●○○○○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○ ○

Distributed Computing and Challenges

What do I Mean by Distributed Computing?

A system where we can take a task and break it down into multiple parts, where each part is worked upon individually.

Challenges of Distributed Computing

Distributed Computing platforms experience these challenges:

Scalability As the network grows, more resources are spent on maintaining and organizing the network.

Georgia State University

Challenges of Distributed Computing

Distributed Computing platforms experience these challenges:

Scalability As the network grows, more resources are spent on maintaining and organizing the network.

Fault-Tolerance As more machines join the network, there is an increased risk of failure.

Georgia State University

Challenges of Distributed Computing

Distributed Computing platforms experience these challenges:

Scalability As the network grows, more resources are spent on maintaining and organizing the network.

Fault-Tolerance As more machines join the network, there is an increased risk of failure.

Load-Balancing Tasks need to be evenly distributed among all the workers.

Georgia State University

Introduction Background Previous Work · ○
○○○●○○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○

UrDHT Autonomous Load-Balancing
○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○ ○

Conclusion

What Are Distributed Hash Tables?

Distributed Key/Value Stores

Distributed Hash Tables are mechanisms for storing values associated with certain keys.

- Values, such as filenames, data, or IP/port combinations are associated with keys.
- These keys are generated by taking the hash of the value.
- We can get the value for a certain key by asking any node in the network.

Georgia State
University

Introduction Background Previous Work UrDHT Autonomous Load-Balancing Conclusion
○○○○●○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○ ○

What Are Distributed Hash Tables?

How Does It Work?

- DHTs organize a set of nodes, each identified by an **ID**.
- Nodes are responsible for the keys that are closest to their IDs.
- Nodes maintain a small list of other peers in the network.
  - Typically a size $\log(n)$ subset of all nodes in the network.
- Each node uses a very simple routing algorithm to find a node responsible for any given key.

Georgia State
University

Current Applications

Applications that use or incorporate DHTs:

- P2P File Sharing applications, such as BitTorrent.
- Distributed File Storage.
- Distributed Machine Learning.
- Name resolution in a large distributed database.

Georgia State
University

Strengths of DHTs

DHTs are designed for large P2P applications, which means they need to be (and are):

- Scalable
- Fault-Tolerant
- Load-Balancing

DHTs Address the Specified Challenges

The big issues in distributed computing can be solved by the
mechanisms provided by Distributed Hash Tables.

Georgia State
University

Uses For DHT Distributed Computing

The generic framework we are proposing would be ideal for:

- Embarrassingly Parallel Computations
  - Any problem that can be framed using Map and Reduce.
  - Brute force cryptography.
  - Genetic algorithms.
  - Markov chain Monte Carlo methods.

- Use in either a P2P context or a more traditional deployment.

Georgia State
University

Table of Contents

Georgia State
University

Required Attributes of DHT

- A distance and midpoint function.
- A closeness or ownership definition.
- A Peer management strategy.

Introduction   Background   Previous Work   UrDHT   Autonomous Load-Balancing   Conclusion
ooooooooo o●ooooooo   oooooooooooooooooo o   ooooooooooooooooooooooooooooooooooooo o

The Components and Terminology

Chord's Closest Metric.

Figure: A Voronoi diagram for a Chord network, using Chord's definition of closest.

Chord Using A Different Closest Metric

Figure: A Voronoi diagram for a Chord network, where closest is defined by the node being the closest in either direction.

Introduction   Background   Previous Work   UrDHT   Autonomous Load-Balancing   Conclusion
○○○○○○○○○   ○○○●○○○○○   ○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○ ○

The Components and Terminology

Terms and Variables

- Network size is n nodes.
- Keys and IDs are m bit hashes, usually SHA1.
- Peerlists are made up of:

  Short Peers   The neighboring nodes that define the network's topology.

  Long Peers   Routing shortcuts.

Georgia State
University

Functions

lookup(*key*) Finds the node responsible for a given key.

put(*key*, *value*) Stores *value* at the node responsible for *key*, where *key* = *hash*(*value*).

get(*key*) Returns the *value* associated with *key*.

Georgia State
University

Chord

- Ring Topology
- Short Peers: predecessor and successor in the ring.
- Responsible for keys between their predecessor and themselves.
- Long Peers: $\log n$ nodes, where the node at index i in the peerlist is

$$root(r + 2^{i-1} \mod m), 1 < i < m$$

Georgia State
University

A Chord Network

Figure: An 8-node Chord ring where $m = 8$. Node 24's long peers are shown.

Fault Tolerence in Chord

- Local maintenance thread gradually fixes the network topology.
  - Each node "notifies" its successor.
  - The successor replies with a better successor if one exists.
- The long peers are gradually updated by performing a lookup on each entry.

Georgia State
University

Handling Churn in General

- Short peers, the neighbors, are regularly queried to:
  - See of the node is still alive.
  - See if the neighbor knows about better nodes.
- Long peer failures are replaced by periodic maintenance.

Table of Contents

Georgia State
University

Overarching Goal

My research has been focused on:

- Abstracting out DHTs.
- Distributed computation using DHTs.

Georgia State
University

Goals

VHash sprung from two related ideas:

- We wanted a way be able optimize latency by embedding it into the routing overlay.

Georgia State University

Goals

VHash sprung from two related ideas:

- We wanted a way be able optimize latency by embedding it into the routing overlay.
- We wanted to create a DHT based off of Voronoi tessellations. Unfortunately:

Goals

VHash sprung from two related ideas:

- We wanted a way be able optimize latency by embedding it into the routing overlay.
- We wanted to create a DHT based off of Voronoi tessellations. Unfortunately:
  - Distributed algorithms for this problem don't really exist.

Georgia State
University

Goals

VHash sprung from two related ideas:

- We wanted a way be able optimize latency by embedding it into the routing overlay.
- We wanted to create a DHT based off of Voronoi tessellations. Unfortunately:
  - Distributed algorithms for this problem don't really exist.
 - Existing approximation algorithms were unsuitable.

Georgia State
University

Introduction   Background   **Previous Work**   UrDHT   Autonomous Load-Balancing   Conclusion
○○○○○○○○○   ○○○○○○○○○   ○●○○○○○○○○○○○○○○○○○   ○

VHash

Voronoi Tesselation

Delaunay Triangulation

DHT and Voronoi Relationship

- We can view DHTs in terms of Voronoi tessellation and Delaunay triangulation.
    - The set of keys the node is responsible for is its Voronoi region.
    - The nodes neighbors are its Delaunay neighbors.

Georgia State University

Introduction Background **Previous Work** UrDHT Autonomous Load-Balancing Conclusion
000000000 000000000 0000●0000000000000 0 0000000000000000000000000000000000000 0

VHash

VHash

- Voronoi-based Distributed Hash Table based on this relationship.
- Uses our approximation to solve for Delaunay neighbors, called DGVH.
- Topology updates occur via gossip-based protocol.
- Routing speed is $O(\sqrt[d]{n})$
- Memory Cost
  - Worst case: $O(n)$
  - Expected maximum size: $\Theta(\frac{\log n}{\log \log n})$

Georgia State
University

Distributed Greedy Voronoi Heuristic

- Assumption: The majority of Delaunay links cross the corresponding Voronoi edges.
- We can test if the midpoint between two potentially connecting nodes is on the edge of the Voronoi region.
- This intuition fails if the midpoint between two nodes does not fall on their Voronoi edge.

Georgia State
University

Introduction Background **Previous Work** UrDHT Autonomous Load-Balancing Conclusion
○○○○○○○○○○ ○○○○○○○○○ ○○○○○○●○○○○○○○○○○○ ○ ○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○ ○

VHash

DGVH Heuristic

   1: Given node n and its list of *candidates*.

   2: *peers* $\leftarrow$ empty set that will contain n's one-hop peers

   3: Sort *candidates* in ascending order by each node's distance to n

   4: Remove the first member of *candidates* and add it to *peers*

   5: **for all** c in *candidates* **do**

   6:    m is the midpoint between n and c

   7:    **if** Any node in *peers* is closer to m than n **then**

   8:       Reject c as a peer

   9:    **else**

 10:       Remove c from *candidates*

 11:       Add c to *peers*

 12:    **end if**

 13: **end for**

Georgia State University

DGVH Time Complexity

For k candidates, the cost is:

$$k \cdot \lg(k) + k \text{ midpoints} + k^2 \text{ distances}$$

However, the expected maximum for k is $\Theta(\frac{\log n}{\log \log n})$, which gives an expected maximum cost of

$$O(\frac{\log^2 n}{\log^2 \log n})$$

Georgia State
University

Results

Figure: As the size of the graph increases, we see approximately 1 error per node.

Results

Figure: These figures show, starting from a randomized network, VHash forms a stable and consistent network topology.
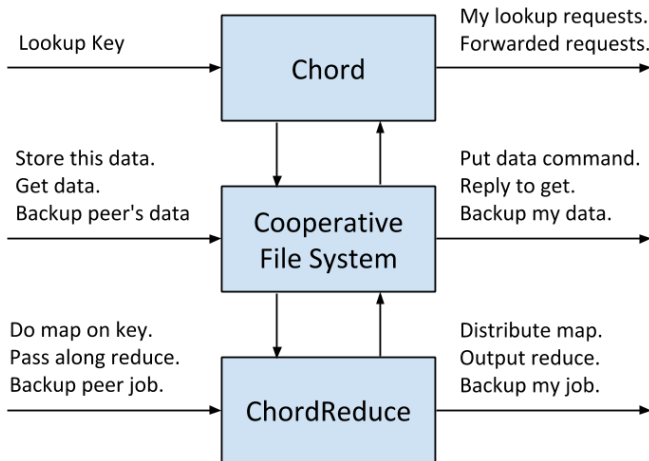
Results

Figure: Comparing the routing effectiveness of Chord and VHash.

Conclusions

- DGVH is simple approximation for Delaunay Triangulation that guarantees a fully connected graph.
- VHash can optimize over a metric such as latency and achieve superior routing speeds as a result.

Georgia State
University

Goals

- We wanted build a more abstract system for MapReduce.
- We remove core assumptions:
    - The system is centralized.
    - Processing occurs in a static network.
- The resulting system must be:
    - Completely decentralized.
    - Scalable.
    - Fault tolerant.
    - Load Balancing.

Georgia State University

Introduction   Background   **Previous Work**   UrDHT   Autonomous Load-Balancing   Conclusion
○○○○○○○○○○ ○○○○○○○○ ○○○○○○○○○○○○○○○●○○○○ ○ ○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○ ○

ChordReduce

System Architecture

Lookup Key → **Chord** → My lookup requests. Forwarded requests.

Store this data. Get data. Backup peer's data → **Cooperative File System** → Put data command. Reply to get. Backup my data.

Do map on key. Pass along reduce. Backup peer job. → **ChordReduce** → Distribute map. Output reduce. Backup my job.

Mapping Data

Figure: The stager sends a map task for each key in the keyfile. In larger networks, this process is streamlined by recursively bundling keys and sending them to the best finger.

ChordReduce

Reducing Results of Data

Figure: Results are sent back via the overlay. If a node receives multiple results, they are reduced before being sent on.

Churn Results

Churn rate per second	Average runtime (s)	Speedup vs 0% churn
0.8%	191.25	2.15
0.4%	329.20	1.25
0.025%	431.86	0.95
0.00775%	445.47	0.92
0.00250%	331.80	1.24
0%	441.57	1.00

Table: The results of calculating π by generating 10^8 samples under churn. Churn is the chance for each node to join or leave the network. The large speedup is from joining nodes acquiring work during experimental runtime.

Georgia State
University

Conclusions

Our experiments established:

- ChordReduce can operate under high rates of churn.
- Execution follows the desired speedup.
- Speedup occurs on sufficiently large problem sizes.

This makes ChordReduce an excellent platform for distributed and concurrent programming in cloud and loosely coupled environments.

Georgia State University

Table of Contents

Georgia State
University

Introduction

- Build on DGVH and VHash

Table of Contents

Georgia State
University

Introduction

- In this project, we set out to confirm the results of ChordReduce
- Objectives:
  - Confirm that high levels of churn can help a DHT based computing environment.
  - Develop better strategies than randomness

Georgia State
University

Distribution of Keys in A DHT

Distribution in Networks of Different Sizes

Table: The median distribution of tasks (or files) among nodes. We can see the standard deviation is fairly close to the expected mean workload ($\frac{tasks}{nodes}$). Each row is the average of 100 trials. Experiments show there is practically little deviation in the median load of the network.

Nodes	Tasks	Median Workload	σ
1000	100000	69.410	137.27
1000	500000	346.570	499.169
1000	1000000	692.300	996.982
5000	100000	13.810	20.477
5000	500000	69.280	100.344
5000	1000000	138.360	200.564
10000	100000	7.000	10.492
10000	500000	34.550	50.366
10000	1000000	69.180	100.319

Georgia State University

Distribution of Work in A DHT

Figure: The probability distribution of workload in a DHT with 1000 nodes and 1,000,000 tasks or files. The vertical dashed line designates the median.

Distribution of Work in Chord

Figure: A visual example of data and nodes in a Chord DHT with 10 nodes (represented by red circles) and 100 tasks (blue pluses).

Terms and Assumptions

- Time is measured in ticks.
- A tick is enough time to perform aggressive, reactive maintenance[1]
- Jobs are measured in tasks; each task can correspond to a file or a piece of a file

---

[1]This has been implemented and tested.

Variables

- Strategy
- Homogeneity
- Work Measurement
- Number of Nodes
- Number of Tasks
- Churn Rate
- Max Sybils or Node Strength
- Sybil Threshold
- Number of Successors

Georgia State University

Introduction Background Previous Work   UrDHT **Autonomous Load-Balancing**   Conclusion
○○○○○○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○●○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○ ○

Experimental Setup

Output

- Ideal Runtime

- Runtime

- Runtime Factor

- Task Distribution

Introduction Background Previous Work   UrDHT **Autonomous Load-Balancing**   Conclusion
○○○○○○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○○●○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○ ○

Churn

Strategy

- The network load-balances using churn
- `churnRate` chance per tick for each node to leave network
- Pool of potentially joining joins at the same rate

Georgia State
University

Runtime

Table: Runtime factor of networks of varying sizes and number of tasks, each using the Churn strategy to load-balance. Each result is the average of 100 trials. The networks are homogeneous and each node consumes one task per tick. A runtime of 1 is the ideal and target.

Churn Rate	10^3 nodes, 10^5 tasks	10^3 nodes, 10^6 tasks	100 nodes, 10^4 tasks	100 nodes, 10^5 tasks	100 nodes, 10^6 tasks
0	7.476	7.467	5.043	5.022	5.016
0.0001	7.122	5.732	4.934	4.362	3.077
0.001	6.047	3.674	4.391	3.019	1.863
0.01	3.721	2.104	3.076	1.873	1.309

Georgia State University

Churn vs Runtime factor

Figure: This graph shows the effect churn has on runtime in a distributed computation. Runtime is measured as how many times slower the computation runs than an ideal computation, where each node receives an equal number of tasks. The lower the runtime factor, the closer it is

Introduction Background Previous Work   UrDHT **Autonomous Load-Balancing**   Conclusion
○○○○○○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○○○○●○○○○○○○○○○○○○○○○○○○○○○○○○○○○ ○

Churn

Churn vs Average Work Per Tick

Figure: With more and more churn in the network, new nodes have a higher chance of joining the network and acquiring work from overloaded nodes. This results in more average work being done each tick, as there are less nodes simply idling.

Introduction Background Previous Work   UrDHT **Autonomous Load-Balancing**   Conclusion
○○○○○○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○○○○○●○○○○○○○○○○○○○○○○○○○○○○○○○ ○

Churn

Base Work Distribution

Figure: The initial distribution of the workload in both networks. As both
networks start with same initial configuration, the distribution is currently
identical. This greatly resembles the distribution we saw in Figure 9.

Introduction   Background   Previous Work   UrDHT   Autonomous Load-Balancing   Conclusion
○○○○○○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○○○○○●○○○○○○○○○○○○○○○○○○○○○○○ ○

Churn

Churn Distribution after 5 Ticks

Figure: This is the distribution of the workloads at the beginning of tick 5. We can see the network using 0.01 churn has fewer nodes with less work and more nodes with a greater workload.

Churn Distribution after 35 Ticks

Figure: After 35 ticks, the effects of churn on the workload distribution become more pronounced. More nodes have consumed all their tasks and are simply idling, but significantly less in the network using churn.

Georgia State University

Remarks

- Diminishing returns
- Maintenance costs can get excessive
- We don't actually have to kill nodes, most of the speedup is from joining.

Introduction Background Previous Work   UrDHT **Autonomous Load-Balancing**   Conclusion
○○○○○○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○○○○○○○○○○●○○○○○○○○○○○○○○○○○○○○○○ ○

Random Injection

Strategy

- Nodes with loads $\leq$ `sybilThreshold` create Sybils
- This check occurs every 5 ticks before work is performed
- These Sybils are randomly placed
- Act as a virtual node so the same node essentially exists in multiple locations
- Sybils are removed if the node that created it has no work

Georgia State
University

Effects of Network Size

- A homogeneous, 1000 node/100,000 task network, never have an average runtime factor greater than 1.7
- Minimum was 1.36.
- In the same network with 1,000,000 tasks, these runtimes were 1.25 and 1.12 respectively.
- On average, the 1,000,000 task network had a runtime factor 0.82 less than the 100,000 task network.

Georgia State
University

Random injection vs no strategy after 5 ticks
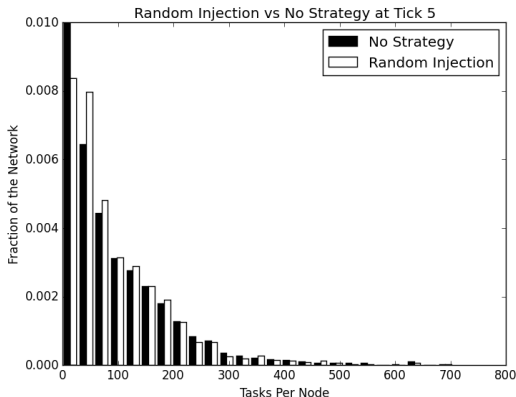
Figure: The networks after 5 ticks. The network using random injection has significantly less underutilized nodes, even after only 5 ticks.

Introduction Background Previous Work   UrDHT **Autonomous Load-Balancing**   Conclusion
○○○○○○○○○ ○○○○○○○○ ○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○○○○○○○○○○○○●○○○○○○○○○○○○○○○ ○

Random Injection

Random Injection vs No Strategy After 35 ticks

Figure: The networks after 35 ticks. The network using random injection has significantly less underutilized nodes and substantially more notes with some or lots of work.

Introduction   Background   Previous Work   UrDHT   **Autonomous Load-Balancing**   Conclusion
oooooooooo ooooooooooo oooooooooooooooooo o   oooooooooooooooooo**o**ooooooooooooooooo o

Random Injection

Random Injection in a Heterogeneous Network

Figure: The workload distribution of heterogenous networks after 35 ticks.

Random Injection VS Churn

Figure: The networks after 35 ticks. The network using random injection load-balances significantly better than the network using Churn.

Impacts of Variables

- `sybilThreshold` would lower the runtime factor.

- Churn had no significant effect

- `Maxsybils` (node strength) had no effect in homogeneous networks

Remarks

- Best runtime factor
- Could still incur high maintenance costs, especially with nodes being deleted as soon as they are made

Georgia State
University

Strategy

- Rather than creating Sybils randomly, nodes create one in their successors
- Finding node id uses mashes
- Estimates which successor has most work.
- Tested estimation against smart method.

Georgia State
University

Neighbor Injection after 5 Ticks

Figure: While the network no strategy appears to have less nodes idling, the active nodes in the network using neighbors are being better utilized.

Neighbor Injection after 35 Ticks

Figure: Despite have more idling nodes, we see that the nodes using the neighbor injection strategy have acquired smaller workloads and have effectively shifted part of the histogram left.

Introduction Background Previous Work   UrDHT **Autonomous Load-Balancing**   Conclusion
○○○○○○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○○○○○○○○○○○○○○○○○○○○●○○○○○○○○○ ○

Neighbor Injection

Smart Neighbor Injection after 5 Ticks

Figure: The workload in the networks after 5 ticks. We can that the network using neighbor injection has directed more nodes that would be idling to perform work. This is especially apparent in comparison to Figure 20.

Smart Neighbor Injection after 35 Ticks

Figure: After 35 ticks, we see the network using the smart neighbor injection strategy has significantly less nodes with little or no work, more nodes with smaller amounts of work, and less nodes with large amounts of tasks.

Impacts of Variables

-

Remarks

-

Strategy

- Nodes with too much work ask for help
- Predecessor with smallest workload and below `sybilThreshold` creates a Sybil
- Reactive vs Proactive

Invitation at 5 Ticks


Figure: Here, we see that the invitation strategy performs markedly better at distributing the work than no strategy. At tick 5, there are many more nodes with small work loads. There are also few nodes with large amounts of work.

Invitation at 35 Ticks

Figure: At 35 ticks, we can see the network using the invitation strategy perform markedly better than the network using no strategy. The highest load is around 500 tasks in the network using invitation, compared to approximately 650 tasks in the network using no strategy.

Invitation vs Smart Neighbors at Tick 5

Figure: A comparison of a network using the smart neighbors strategy and a network using the invitation strategy. Invitation has slightly more nodes on the performing less work, but no significant differences have emerged.

Introduction Background Previous Work UrDHT Autonomous Load-Balancing Conclusion
○○○○○○○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○ ○ ○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○●○○ ○

Invitation

Invitation vs Smart Neighbors at Tick 35

Figure: After 35 ticks, differences between the two strategies have emerged. The network using the invitation strategy has significantly less nodes with a small work load and many more with large work loads.

Remarks

-

Summary

- Reactive vs Proactive
- Heterogeneity was best handled by Churn or Random Injection.

Table of Contents

Georgia State
University

Introduction Background Previous Work   UrDHT Autonomous Load-Balancing   Conclusion
○○○○○○○○○ ○○○○○○○○○ ○○○○○○○○○○○○○○○○○○○○ ○   ○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○ ●

Conclusion

Conclusion

- Load Balancing